

# Cracking bcrypt on FPGAs

04. June 2014

FluxFingers  
Ruhr University Bochum  
Friedrich Wiemer

RUB



## Cracking bcrypt on FPGAs

2014-06-04



# Outline

1 Motivation

2 FPGAs

3 bcrypt

4 Design of Cracker

5 Demo

6 Results

2014-06-04

└ Outline

# Motivation

Why do we care?

- password cracking has an inherent parallel structure
- FPGAs enable to exploit true parallelism
- bcrypt claims to resists hardware optimisations
- implementation presented last year<sup>1</sup> was... not optimal :-)
- play with expensive hardware

<sup>1</sup> <http://www.openwall.com/presentations/Passwords13-Energy-Efficient-Cracking/>

2014-06-04

## Cracking bcrypt on FPGAs

### └ Motivation

#### └ Motivation

- OpenWall bcrypt implementation on different platforms:
  - (mainly) parallela/epiphany board
  - zedboard
  - Xeon Phi
  - Haswell
- Problems with FPGA implementation:
  - only implemented cost loop in fabric
  - wait cycles
  - 14 cores with very unbalanced resource utilization



2014-06-04

## Cracking bcrypt on FPGAs

- └ FPGAs

What are FPGAs?

short: (re-)programmable hardware

long: see next slides



# What are FPGAs?

GPCPUs ↔ FPGAs ↔ ASICs

Field Programmable Gate Arrays lie between General Purpose CPUs and Application Specific Integrated Circuits. They combine advantages from both.

## Advantages

- fast to program
- low level hardware design
- re-programmable

## Drawbacks

- slower clock speed
- lower resource utilization

In contrast to CPUs, parallelism is directly expressible, but one does not have to design everything down to transistor level.

2014-06-04

## Cracking bcrypt on FPGAs

### FPGAs

#### What are FPGAs?

- CPUs have general functions
- supports wide range of applications  
⇒ good in many points but not best in specific applications
- ASICs require huge design effort
- tedious modeling and simulation down to gate-level  
⇒ fastest hardware for specific applications
- FPGAs easier to program than ASICs
- re-programmable (loadable with task specific programs)
- allow low level hardware design while still abstracts from gate level

RUHR-UNIVERSITÄT BOCHUM  
What are FPGAs?  
GPCPUs ↔ FPGAs ↔ ASICs

Field Programmable Gate Arrays lie between General Purpose CPUs and Application Specific Integrated Circuits. They combine advantages from both.

Advantages	Drawbacks
■ fast to program ■ low level hardware design ■ re-programmable	■ slower clock speed ■ lower resource utilization

In contrast to CPUs, parallelism is directly expressible, but one does not have to design everything down to transistor level.

# What are FPGAs?

How does it work?

FPGAs are build mainly out of

- Look Up Tables (LUTs)
- Flip Flops (FFs)
- Block RAMs (BRAMs)
- Digital Signal Processing units (DSPs)



but can contain hard instantiated cores  
(e.g. ARM Cortex A9, etc.)

## Cracking bcrypt on FPGAs

### └ FPGAs

#### └ What are FPGAs?

contains different kind of logic blocks

- LUTs are needed for custom logic function
- FFs to register LUT outputs, build local storage, reduce critical path, re-programmable routes
- BRAMs bigger local storage, still accessible within one clock cycle
- DSPs for signal processing (multiply accumulate)
- other hard instantiated cores (e.g. ARM Cortex A9)



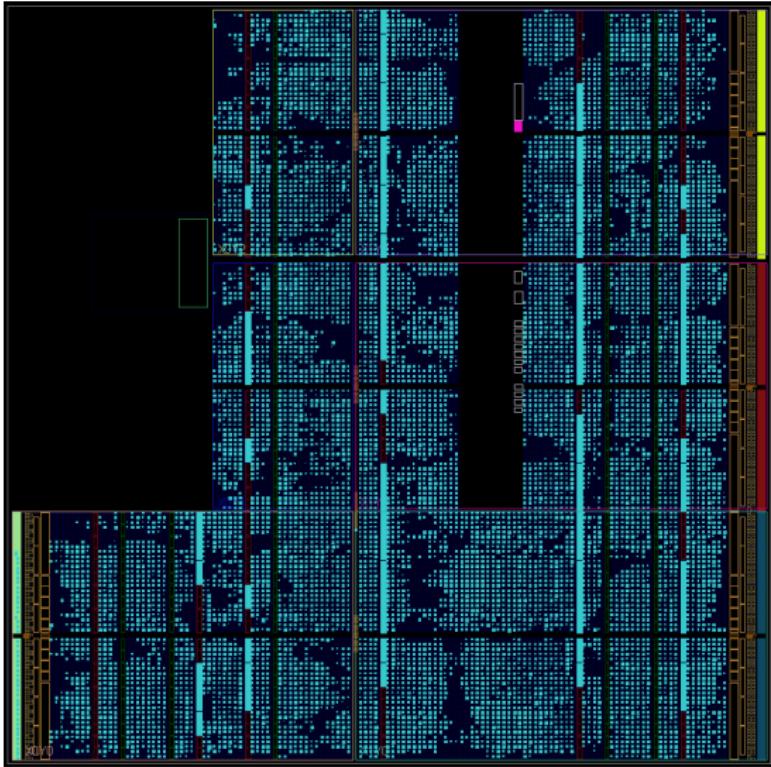
FPGAs are build mainly out of

- Look Up Tables (LUTs)
- Flip Flops (FFs)
- Block RAMs (BRAMs)
- Digital Signal Processing units (DSPs)

but can contain hard instantiated cores  
(e.g. ARM Cortex A9, etc.)

# What are FPGAs?

Artix7

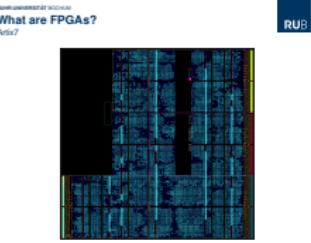


2014-06-04

## Cracking bcrypt on FPGAs

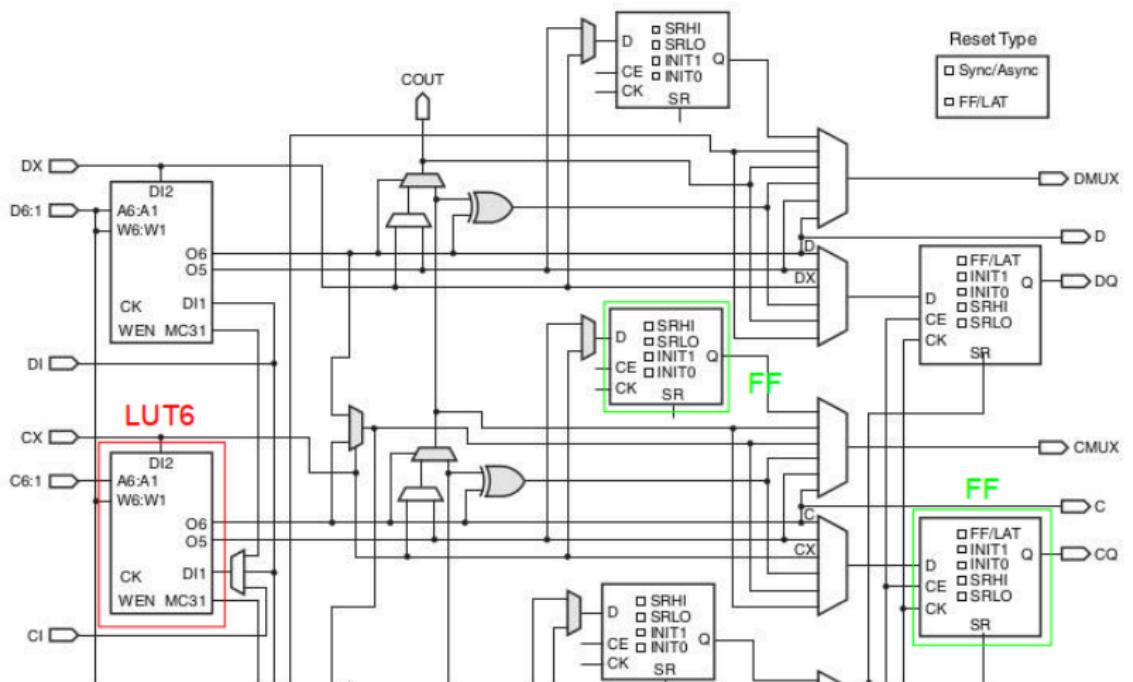
- └ FPGAs

- └ What are FPGAs?



# What are FPGAs

## Slice



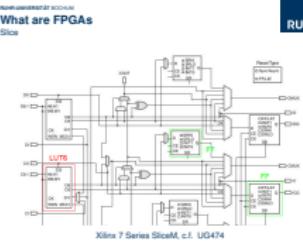
Xilinx 7 Series SliceM, c.f. UG474

## Cracking bcrypt on FPGAs

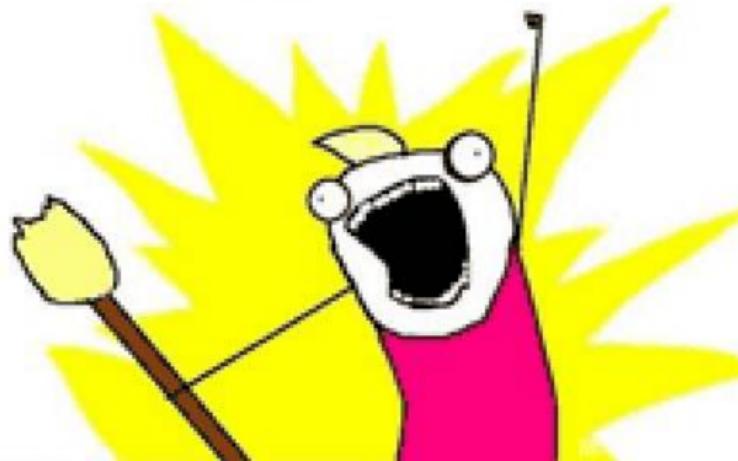
### FPGAs

#### What are FPGAs

- 4 LUT6, 8 FFs per slice (2 FFs per LUT6)
- additional muxer, fast carry path, etc.
- is this important to know?
- well.. depends :-)
- one has to know how the hardware is build up, to get best performance (respect the hardware)



# CLUSTER ALL THE FPGAs!!!

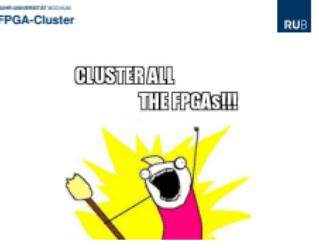


9

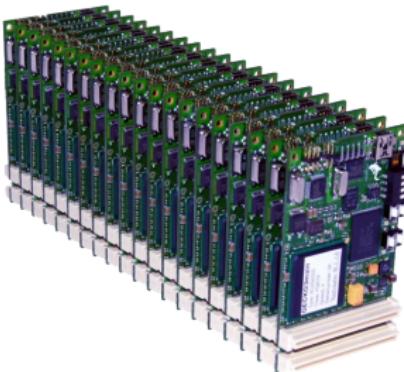
2014-06-04 Cracking bcrypt on FPGAs  
└ FPGAs  
    └ FPGA-Cluster

What are FPGA Cluster?

- similar to GPU cluster
- aim: improve parallel performance



FPGA Cluster like  
COPACOBANA or RIVYERA  
combine many FPGA chips.



### *Can we use this?*

- Yes!
- split password range between every FPGA
- with small overhead, this results in a nice speed up

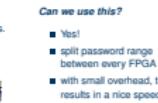
2014-06-04

## Cracking bcrypt on FPGAs

### └ FPGAs

#### └ FPGA-Cluster

- example for cluster:
  - COPACOBANA
    - 120 low power Spartan3 FPGAs
  - RIVYERA
    - 64 Spartan6 FPGAs
- Why is this useful for us?
  - Password Range can be split up at synthesis time
  - every FPGA works independently
  - with few dependencies, this results in quasi linear speed up



- Can we use this?
- Yes!
  - split password range between every FPGA
  - with small overhead, this results in a nice speed up

# What is bcrypt?

Introduced in 1999 by Provos and Mazières.<sup>2</sup>

## *EksBlowfish*

- block cipher based on blowfish
- expensive key schedule
- encryption equals blowfish encryption

Implemented in OpenBSD 2.1, Ruby on Rails, PHP as standard password hash.

<sup>2</sup> [http://www.usenix.org/events/usenix99/full\\_papers/provos/provos.pdf](http://www.usenix.org/events/usenix99/full_papers/provos/provos.pdf)

2014-06-04

## Cracking bcrypt on FPGAs

- └ bcrypt

### └ What is bcrypt?

- proposed two constructions
- *EksBlowfish* less known, but serves as underlying building block for bcrypt
- itself is based on blowfish, expensive key schedule
- *bcrypt* cost parameterized hash function
- cost factor controls computational efford of *EksBlowfishSetup*
- practical applications in OpenBSD, Ruby on Rails, PHP
- We start with explaining the key schedule

What is bcrypt?	
Introduced in 1999 by Provos and Mazières. <sup>2</sup>	
<b>EksBlowfish</b>	<b>bcrypt</b>
<ul style="list-style-type: none"> <li>■ block cipher based on blowfish</li> <li>■ expensive key schedule</li> <li>■ encryption equals blowfish encryption</li> </ul>	<ul style="list-style-type: none"> <li>■ cost-parameterized</li> <li>■ uses <i>EksBlowfishSetup</i> (key schedule) and encryption</li> </ul>
Implemented in OpenBSD 2.1, Ruby on Rails, PHP as standard password hash.	
<small><a href="http://www.usenix.org/events/usenix99/full_papers/provos/provos.pdf">http://www.usenix.org/events/usenix99/full_papers/provos/provos.pdf</a></small>	

# What is bcrypt?

EksBlowfishSetup

## Pseudocode

### Algorithm 1: EksBlowfishSetup

**Input:** cost, salt, key

**Output:** state

*state*  $\leftarrow$  InitState();

*state*  $\leftarrow$  ExpandKey(*state*, *salt*, *key*);

**Repeat** ( $2^{cost}$ ) **begin**

*state*  $\leftarrow$  ExpandKey(*state*, 0, *salt*);

*state*  $\leftarrow$  ExpandKey(*state*, 0, *key*);

**end**

**return** *state*;

## Cracking bcrypt on FPGAs

└ bcrypt

└ What is bcrypt?

2014-06-04

- state consists of 4 SBoxes, each  $256 \times 32$  bit and
- 18 32 bit P-values (subkeys)  $\Rightarrow$  4 kB memory
- initialized with digits of  $\pi$
- ExpandKey updates state (see next slide)
- number of iterations exponential in cost parameter
- *ExpandKey(state, 0, .)* equals blowfish key schedule

**Pseudocode**  
Algorithm 1: EksBlowfishSetup  
Input: cost, salt, key  
Output: state  
*state*  $\leftarrow$  InitState();  
*state*  $\leftarrow$  ExpandKey(*state*, *salt*, *key*);  
Repeat ( $2^{cost}$ ) begin  
| *state*  $\leftarrow$  ExpandKey(*state*, 0, *salt*);  
| *state*  $\leftarrow$  ExpandKey(*state*, 0, *key*);  
end  
return *state*;

# What is bcrypt?

ExpandKey

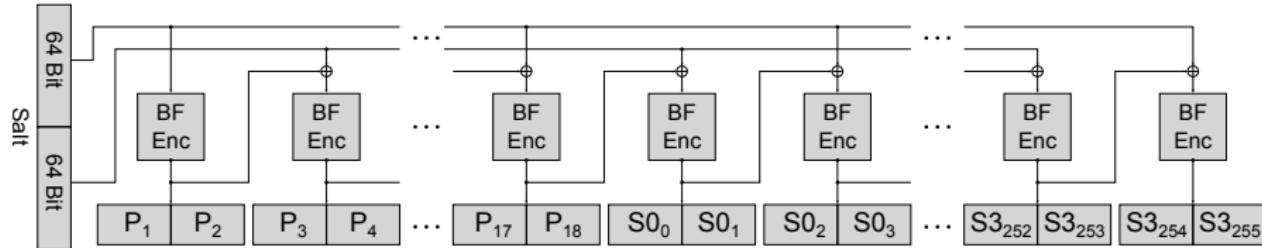


## ExpandKey

Modified version of blowfish key schedule. XORes the key onto subkeys and afterwards, updates the state sequentially.

Calls `bf_enc` 521 times per invocation.

## State Update

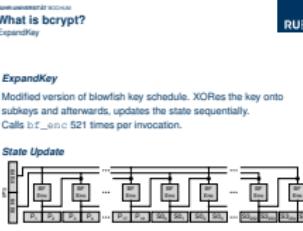


# Cracking bcrypt on FPGAs

## bcrypt

### What is bcrypt?

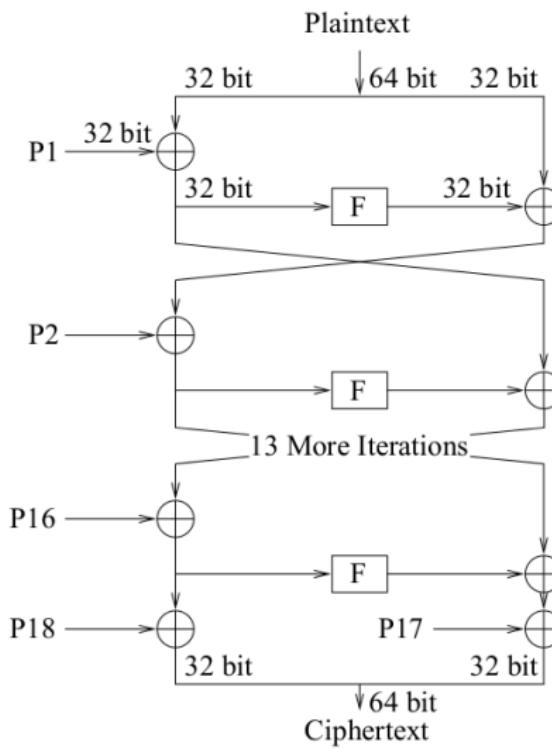
2014-06-04



- ExpandKey blowfish encrypts lower 64 bit of Salt
- Result substitutes  $P_1$ ,  $P_2$  and
- gets xored with higher 64 bit of Salt, encrypted again...
- Alltogether 521 blowfish encryptions per call
- How does blowfish encryption works?

# What is bcrypt?

blowfish encryption

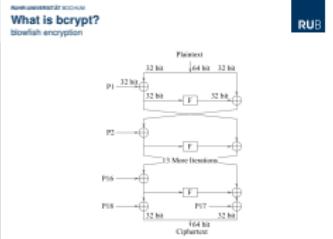


## Cracking bcrypt on FPGAs

### bcrypt

#### What is bcrypt?

- simple 16 round Feistel network
  - in every round: key xor, f-Function xor, swap
  - f-Function consists of 4 SBox look ups, 2 additions, 1 xor
  - finalize: xor  $P_{17}, P_{18}$
  - 64 bit block size
- ⇒ nice minimalistic structure



# What is bcrypt?

bcrypt

RUB

## Pseudocode

### Algorithm 2: bcrypt

**Input:** cost, salt, key

**Output:** hash

*state*  $\leftarrow$  EksBlowfishSetup(*cost*, *salt*, *key*);

*ctext*  $\leftarrow$  "OrpheanBeholderScryDoubt";

**Repeat (64) begin**

*ctext*  $\leftarrow$  EncryptECB(*state*, *ctext*);

**end**

**return** Concatenate(*cost*, *salt*, *key*);

## Cracking bcrypt on FPGAs

└ bcrypt

    └ What is bcrypt?

2014-06-04

RUHR-UNIVERSITÄT BOCHUM  
What is bcrypt?  
bcrypt

---

Pseudocode

---

```
Algorithm 2: bcrypt
Input: cost, salt, key
Output: hash
state  $\leftarrow$  EksBlowfishSetup(cost, salt, key);
ctext  $\leftarrow$  "OrpheanBeholderScryDoubt";
Repeat (64) begin
    | ctext  $\leftarrow$  EncryptECB(state, ctext);
end
return Concatenate(cost, salt, key);
```

- bcrypt itself is quite simple
  - calls EksBlowfishSetup followed by
  - 64 blowfish encryptions in ECB mode of magic value (3 blocks)
- ⇒ almost all work/time is spent during key schedule

## *bcrypt is*

- based on blowfish
- cost-parameterized password hash
- sequential

## *bcrypt needs*

- exponential (in *cost*) loop iterations
- 4 kB memory

# Cracking bcrypt on FPGAs

## └ bcrypt

### └ Where are we?

- bcrypts cost depend only on many sequential blowfish encryption
- authors claim randomly accessed memory is costly

**bcrypt is**

- based on blowfish
- cost-parameterized password hash
- sequential

**bcrypt needs**

- exponential (in *cost*) loop iterations
- 4 kB memory



## Cracking bcrypt on FPGAs

└ Design of Cracker

└ Design of Cracker

2014-06-04



- we have no chance to parallelize bcrypt's structure
- does the memory consumption hurt?
- depends on target FPGAs

# Target Platforms

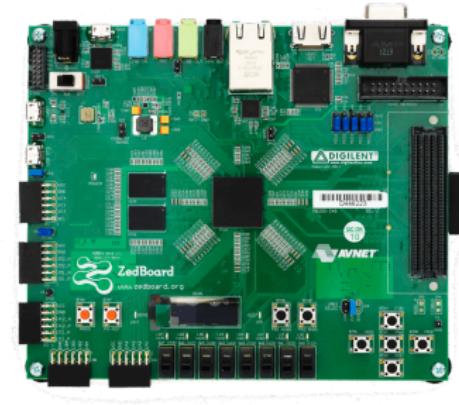
zedboard



Low cost, low power FPGA

## Specs

- Zynq-7000 XC7Z020 FPGA  
(comparable to Artix7)
- ARM Cortex A9 CPU
- HDMI, VGA, Ethernet, Audio,  
USB, JTAG, SD Card,  
Buttons...



2014-06-04

## Cracking bcrypt on FPGAs

### └ Design of Cracker

#### └ Target Platforms

- OpenWall project also targeted zedboard
- combines ARM Cortex A9 with programmable logic
- contains 140 18 kbit BRAMs
- embedded Linux can be run from SD card

RUHR-UNIVERSITÄT BOCHUM  
Target Platforms  
zedboard

Low cost, low power FPGA

Specs

- Zynq-7000 XC7Z020 FPGA  
(comparable to Artix7)
- ARM Cortex A9 CPU
- HDMI, VGA, Ethernet, Audio,  
USB, JTAG, SD Card,  
Buttons...

High Performance FPGA cluster

## Specs

- 64 Spartan6 LX150  
(available with up to 128)
- i7 CPU
- 16 GB RAM



Up to 10 RIVYERAs can be fit into a standard rack.

## Cracking bcrypt on FPGAs

### └ Design of Cracker

### └ Target Platforms

2014-06-04

RUHR-UNIVERSITÄT BOCHUM  
Target Platforms  
RIVYERA

High Performance FPGA cluster  
Specs

- 64 Spartan6 LX150  
(available with up to 128)
- i7 CPU
- 16 GB RAM

Up to 10 RIVYERAs can be fit into a standard rack.



- RIVYERA is the actual FPGA cluster available @ EmSec
- contains up to 128 Spartan6 (ours: 64)
- each Spartan6 268 18 kbit BRAMs  
(roughly twice as big as Artix7 on zedboard, but slower speed grade)

# Low Area Footprint

2014-06-04

Cracking bcrypt on FPGAs

└ Design of Cracker

└ Optimization Goal

Low Area Footprint

Design Goal?

- simple Feistel structure will allow high clock rates
  - ⇒ keep area footprint as low as possible
  - ⇒ fit as many cores onto on FPGA as possible

# Design of Cracker

## API

- keep API minimalistic ⇒ low bandwidth interface
- transfer target *salt* and *hash*
- start cracker
- on success, read *password*

## Password Generation

- on-chip password generation
- needs no bandwidth
- split password range at synthesis-time

## Cracking bcrypt on FPGAs

### Design of Cracker

2014-06-04

### Design of Cracker

#### API

- keep API minimalistic ⇒ low bandwidth interface
- transfer target salt and hash
- start cracker
- on success, read password

#### Password Generation

- on-chip password generation
- needs no bandwidth
- split password range at synthesis-time

## *bcrypt Core*

- memory for SBoxes, subkeys, key and initial values
- almost no logic needed for algorithm
- consists mainly of BRAM look ups

## *Problematic*

- memory for password storage
- registers require too much logic

2014-06-04

## Cracking bcrypt on FPGAs

### └ Design of Cracker

#### └ Design of Cracker

**bcrypt Core**

- memory for SBoxes, subkeys, key and initial values
- almost no logic needed for algorithm
- consists mainly of BRAM look ups

**Problematic**

- memory for password storage
- registers require too much logic

- 2 SBoxes per BRAM, subkeys in 1 extra Bram
  - initial values take another 2 BRAMs (can be shared)
  - but: how to store passwords?
  - registers in logic are too costly
- ⇒ quad-core design

# Design of Cracker

## Quad Core

- bundle four bcrypt cores together
- store four passwords in one BRAM
- during initialization, password generator can access BRAM
- core0, core1 and core2, core3 can access the BRAM alternately

## Advantages

- cores need access only during begin of ExpandKey thus, they only diverge by 36 clock cycles

## Cracking bcrypt on FPGAs

### Design of Cracker

#### Design of Cracker

2014-06-04

##### Quad Core

- bundle four bcrypt cores together
- store four passwords in one BRAM
- during initialization, password generator can access BRAM
- core0, core1 and core2, core3 can access the BRAM alternately

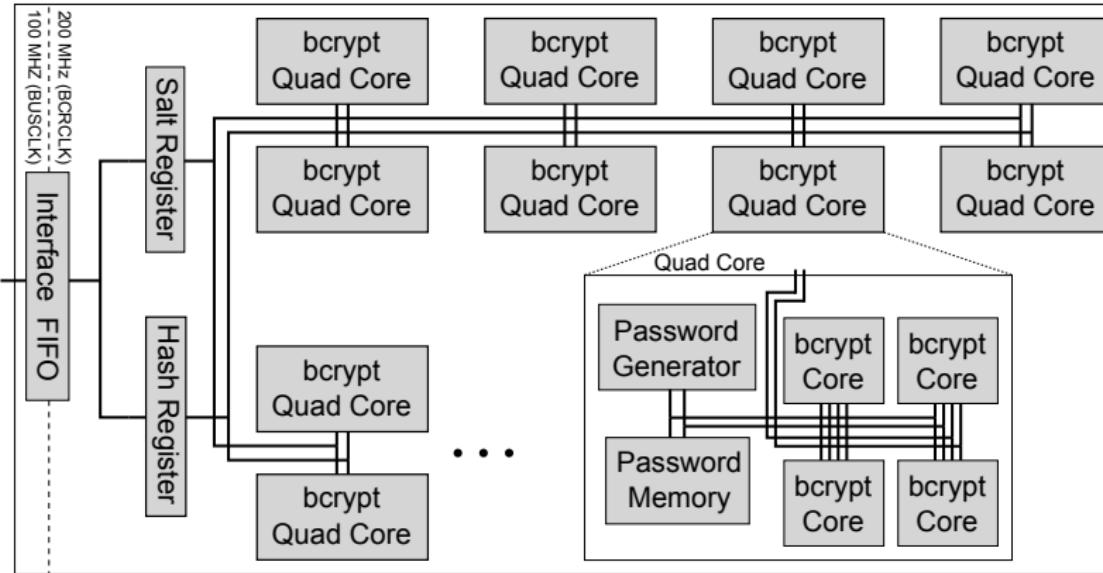
##### Advantages

- cores need access only during begin of ExpandKey thus, they only diverge by 36 clock cycles

- dual port BRAM allows 2 read/write accesses per clock
- bundle 4 cores, alternate their memory access phases
- password generator can access the memory during 256 initialization cycles
- cores diverge only by 36 clock cycles

# Design of Cracker

Resulting Design

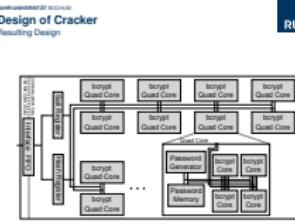


2014-06-04

## Cracking bcrypt on FPGAs

└ Design of Cracker

└ Design of Cracker



- every quad core operates fully independent on his own password subrange

# How to interface the Cracker?

Different possibilities for zedboard:

## ■ AXI4 Interfaces

- + low area footprint
- more complex to implement

## ■ Xillybus

- + abstracts from AXI4 Interfaces
- + Linux support ⇒ easy to demonstrate
- proprietary
- “fat”, i.e. higher resource consumption

Custom interface for RIVYERA needed.

# Cracking bcrypt on FPGAs

## └ Design of Cracker

### └ How to interface the Cracker?

2014-06-04

Different possibilities for zedboard:  
■ AXI4 Interfaces

- + low area footprint
- more complex to implement

  
■ Xillybus

- + abstracts from AXI4 Interfaces
- + Linux support ⇒ easy to demonstrate
- proprietary
- “fat”, i.e. higher resource consumption

  
Custom interface for RIVYERA needed.

# How to interface the Cracker?

## Xillybus

- stream access through FIFOs or DMA into fabric logic
- abstraction from AXI4 is quite nice
- Linux driver ⇒ just write to /dev/xillybus\_mem\_8

## Memory Layout

```
root@localhost:~# hexdump -C -v -n 64 /dev/xillybus_mem_8
00000000  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....
00000010  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....
00000020  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....
00000030  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....
00000040
root@localhost:~#  Salt Hash Password Status Register
```

## Cracking bcrypt on FPGAs

### Design of Cracker

#### └ How to interface the Cracker?

- Linux driver allows DMA into logic
- host side interface consists only of file write/reads.
- Status Register:
  - low nibble used for host-to-FPGA communication (reset, start)
  - high nibble used for FPGA-to-host communication (done, success)

RUHR-UNIVERSITÄT BOCHUM  
How to interface the Cracker?

**Xillybus**

- stream access through FIFOs or DMA into fabric logic
- abstraction from AXI4 is quite nice
- Linux driver ⇒ just write to /dev/xillybus\_mem\_8

**Memory Layout**

```
root@localhost:~# hexdump -C -v -n 64 /dev/xillybus_mem_8
00000000  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....
00000010  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....
00000020  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....
00000030  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....
00000040
root@localhost:~#  Salt Hash Password Status Register
```

# Design of Cracker

Resulting Resources

## **zedboard**

- estimations for one zedboard:  
28 cores as upper bound, BRAMs as limiting resource
- first design attempt (password in registers):  
12 cores fit, LUT utilization way to high
- Quad Core Design:  
24 cores fit, with using “fat” interface

## **RIVYERA**

- Quad Core Design:  
48 cores per Spartan6 ⇒ 3072 cores on whole RIVYERA



2014-06-04

# Cracking bcrypt on FPGAs

## └ Design of Cracker

### └ Design of Cracker

RUHR-UNIVERSITÄT BOCHUM  
Design of Cracker  
Resulting Resources

**zedboard**

- estimations for one zedboard:  
28 cores as upper bound, BRAMs as limiting resource
- first design attempt (password in registers):  
12 cores fit, LUT utilization way to high
- Quad Core Design:  
24 cores fit, with using “fat” interface

**RIVYERA**

- Quad Core Design:  
48 cores per Spartan6 ⇒ 3072 cores on whole RIVYERA



2014-06-04 Cracking bcrypt on FPGAs  
└ Demo  
  └ Demo

- demonstrate interface :-)



# Results

Compared to GPUs

## **cost Parameter and Hash-rates**

		Target (CPU) runtime			
		1ms	10ms	100ms	1000ms
<b>bcrypt cost</b>	3.69	7.03	10.3	13.6	
		Hashs per Second			
		1ms	10ms	100ms	1000ms
<b>zedboard</b>	<b>9,230 H/s</b>	<b>916.25 H/s</b>	<b>98.77 H/s</b>	<b>9.93 H/s</b>	
Spartan-6	9,230 H/s	916.25 H/s	98.77 H/s	9.93 H/s	
GTX 480	2,868 H/s	319.37 H/s	33.73 H/s	2.71 H/s	

## Cracking bcrypt on FPGAs

### Results

#### Results

2014-06-04

- measured bcrypt runtime on different CPUs  
⇒ averaged cost factor
- Hash-rates compared too previous cost factors
- GPU Hash-rate measured with oclHashcat  
⇒ about 3 times faster than GTX 480

cost Parameter and Hash-rates				
Target (CPU) runtime				
	1ms	10ms	100ms	1000ms
bcrypt cost	3.69	7.03	10.3	13.6
Hashs per Second				
	1ms	10ms	100ms	1000ms
<b>zedboard</b>	<b>9,230 H/s</b>	<b>916.25 H/s</b>	<b>98.77 H/s</b>	<b>9.93 H/s</b>
Spartan-6	9,230 H/s	916.25 H/s	98.77 H/s	9.93 H/s
GTX 480	2,868 H/s	319.37 H/s	33.73 H/s	2.71 H/s

# Results



## *Compared to OpenWall*

- OpenWall achieves 780 H/s with cost 5
- zedboard achieves 3,200 H/s with cost 5

## *Besides pure Hash-rate*

- zedboard consumes 5 W, GTX 480 consumes up to 430 W
- zedboard costs \$300–\$400, GTX 480 cost \$500 in 2010

## └ Results

### └ Results

2014-06-04

#### Compared to OpenWall

- OpenWall achieves 780 H/s with cost 5
- zedboard achieves 3,200 H/s with cost 5

#### Besides pure Hash-rate

- zedboard consumes 5 W, GTX 480 consumes up to 430 W
- zedboard costs \$300–\$400, GTX 480 cost \$500 in 2010

# Results

## Optimizations

### *Optimizations left*

- 200 MHz on zedboard, but only 100 MHz on RIVYERA
- low FF/LUT ratio
- unused FFs can be used to buffer signals
  - ⇒ shorter critical path
  - ⇒ higher clock speed
- Optimizing the interface will lead to 28 cores per zedboard

2014-06-04

*Optimizations left*

- 200 MHz on zedboard, but only 100 MHz on RIVYERA
- low FF/LUT ratio
- unused FFs can be used to buffer signals
  - ⇒ shorter critical path
  - ⇒ higher clock speed
- Optimizing the interface will lead to 28 cores per zedboard

# Questions?

Thanks!



2014-06-04

## Cracking bcrypt on FPGAs

- └ Results
- └ Questions?

